

Kazuma Sasaki — Multilingual Audio Data

Speech data creation · Speaker diarization · Audio captioning · Annotation QA

kazmike.s@gmail.com · [linkedin.com/in/kazumasasaki](https://www.linkedin.com/in/kazumasasaki) · Tokyo

Who I am

2+ years in AI training data. Most recently a QA lead on a frontier AI lab's multimodal audio project, where I grew from trainer → pre-completion reviewer → post-completion reviewer → QA lead. I've worked to frontier-lab quality bars, trained annotators, and built review pipelines from scratch — in Japanese and English.

I specialize in the part of the market that is hardest to source: **high-quality audio data** in Japanese and other Asian languages.

What I do

1. Multilingual speech data creation (core service)

Full annotation built from scratch: accurate transcription, word-level speaker diarization, speaker & voice profiling, and emotion / delivery tagging. Production-ready training data for ASR, TTS, and voice agents.

2. Audio captioning (specialty)

Detailed description of every non-speech sound — music, sound effects, ambient, silence — for multimodal models that need to understand audio, not just words. Few suppliers can do this at frontier-lab quality.

3. Data QA (audit + report)

Review of your existing annotation data, returned as a structured report: per-item findings, the rule each one breaks, and the exact fix.

4. Pipeline consulting & gold sets

Design of guidelines, error taxonomy, scoring rubric and reviewer workflow; plus high-accuracy reference sets to benchmark your models or vendors.

Where I'm strongest

- Multilingual coverage — Japanese & Asian languages (the scarce, high-demand gap)
- Speaker diarization — word-level, ~100ms timestamp accuracy
- Audio captioning — music, sound effects, ambient sound, silence
- Speaker & voice profiling — pitch, accent, vocal texture
- Emotion / delivery tagging — conservative, audio-grounded

How we'd start

A free sample on a small batch of your data — either a QA audit of what you have, or a short annotated sample built to your spec. You see the quality first-hand, then decide whether to scope further. No commitment.

How I price

Every project is scoped individually — pricing depends on language mix, turnaround, volume, and how much of the pipeline you need (full creation, QA only, or consulting). Asian-language work is a specialty.

The easiest first step is free: a sample audit or a short annotated sample, so you can judge the quality before any numbers come up.

Let's start with a free sample — send me a few minutes of audio and I'll show you the quality you'd get.